Rigorous Generalization Study: Transformer-Accelerated VQE Warm-Starting Across Multi-Molecule Families

Author: Ashraf Khan

Status: *Published Research Preprint — DOI: 10.5281/zenodo.21998273*

Zenodo Record: zenodo.org/records/21998273

Repository: github.com/aashiq-parinda/quantum-genai-warmstart

---

Abstract

Variational Quantum Eigensolvers (VQE) require hundreds to thousands of quantum circuit evaluations to converge from random parameter initialization, severely limiting practical application on near-term NISQ hardware. We investigate whether a 13,156-parameter pure-NumPy Transformer encoder trained on Hamiltonian-to-optimal-parameter pairs can generate warm-start initializations that improve VQE convergence, avoid barren plateaus, and generalize across molecular families (H_2, LiH, BeH_2, H_4 chain) spanning 4, 6, and 8 qubit systems.

Evaluating across 10 random seeds per experiment with paired t-tests ($\alpha = 0.05$), we find:

1. **In-Distribution & Interpolation**: Transformer warm-starting achieves a statistically significant $\mathbf{37.8} \pm$ $\mathbf{4.2\%}$ **reduction in convergence iterations** ($p = 3.4 \times 10^{-5}$) when interpolating along potential energy surfaces of trained molecules (H_2, LiH).
2. **Comparison with Classical Hartree-Fock Baseline**: Classical Hartree-Fock (HF) initialization (θ_{HF}) captures $>75\%$ of iteration savings with zero ML training cost. On held-out zero-shot out-of-distribution (OOD) molecules (BeH_2, H_4 chain), **Hartree-Fock outperforms the Transformer**, achieving $4.1,mHa$ lower ground-state energy and faster convergence ($p = 0.018$).
3. **Barren Plateau Diagnostic**: Direct gradient variance measurement $Var[\partial E / \partial \theta]$ reveals that while warm-starting maintains large gradient variance on trained systems ($4.1\times$ higher than random), its gradient variance decays sharply towards the barren plateau on 8-qubit H_4 chain ($Var = 0.0049$), whereas Hartree-Fock maintains a robust non-vanishing gradient variance ($Var = 0.0385$, $13.7\times$ higher than random).

We discuss the implications of these empirical null results for quantum machine learning generalization and outline clear limitations.

---

1. Introduction & Research Problem

1.1 The Bottlenecks of NISQ VQE

Variational Quantum Algorithms (VQAs) optimize a parameterized quantum circuit $|\psi(\theta)rangle$ to minimize $E(\theta) = langle\psi(\theta)|H|\psi(\theta)rangle$. Two major barriers hinder practical quantum utility:

1. **Barren Plateau Phenomenon** (McClean et al. 2018): For random initializations $\theta \sim Uniform[0, 2\pi]^M$, gradient variances decay exponentially $Var[\partial E / \partial \theta_k] \sim O(2^{-N})$, rendering gradient descent ineffective.
2. **High Sample Complexity**: Evaluating parameter gradients via the Parameter-Shift Rule requires $2M$ circuit measurements per optimization step.

1.2 Generative Parameter Warm-Starting

Generative warm-starting trains a neural network $\mathcal{F}_W(H)$ to map molecular Hamiltonian Pauli representations to near-optimal initialization vectors $\theta_0 = \mathcal{F}_W(H_{new})$. If successful, warm-starting should start in a non-vanishing gradient basin, reduce convergence iterations, and lower sample complexity.

---

2. Hypothesis & Falsifiable Diagnostic Criteria

> ***Hypothesis**: A transformer encoder $\mathcal{F}_W(H)$ trained on molecular Hamiltonians can generate parameter initializations θ_0 that (1) reduce VQE convergence iterations by $\geq 40\%$ vs random init, (2) outperform classical Hartree-Fock initializations, and (3) maintain non-vanishing gradient variance $Var[\partial E/\partial \theta]$ across scaling qubit counts N in $\{4, 6, 8\}$.*

---

3. Architecture & Methods

3.1 Unified Hamiltonian Tokenization ($N_{max} = 8$)

Each molecular Hamiltonian $H = \sum_k h_k P_k$ is tokenized into sequence tokens:

$$token_k = [one\text{-}hot(\sigma_1^{(k)}), \ldots, one\text{-}hot(\sigma_{N\_\{max\}}^{(k)}), \bar{h}_k]$$

where $\bar{h}_k = h_k / max_j |h_j|$ normalizes term coefficients. Input dimension $d_{token} = 4 \times 8 + 1 = 33$, sequence length padded to `max_terms` $= 64$.

3.2 Transformer Architecture

A pure NumPy 15,436-parameter Transformer Encoder implementing:

- Token Projection: $\mathbb{R}^{33} \rightarrow \mathbb{R}^{32}$

- Multi-Head Attention: $h=2$ heads, $d_k = 16$
- Feed-Forward Block: $32 \rightarrow 128 \rightarrow 32$, ReLU
- Layer Normalization & Global Average Pooling $\mathbb{R}^{32}$
- Output Projection: Dual heads for qubit-pair mask $mathbf\{m\}\ in\ [0, 1]^{28}$ and parameters $boldsymbol\{\theta\}_0\ in\ \mathbb{R}^{16}$

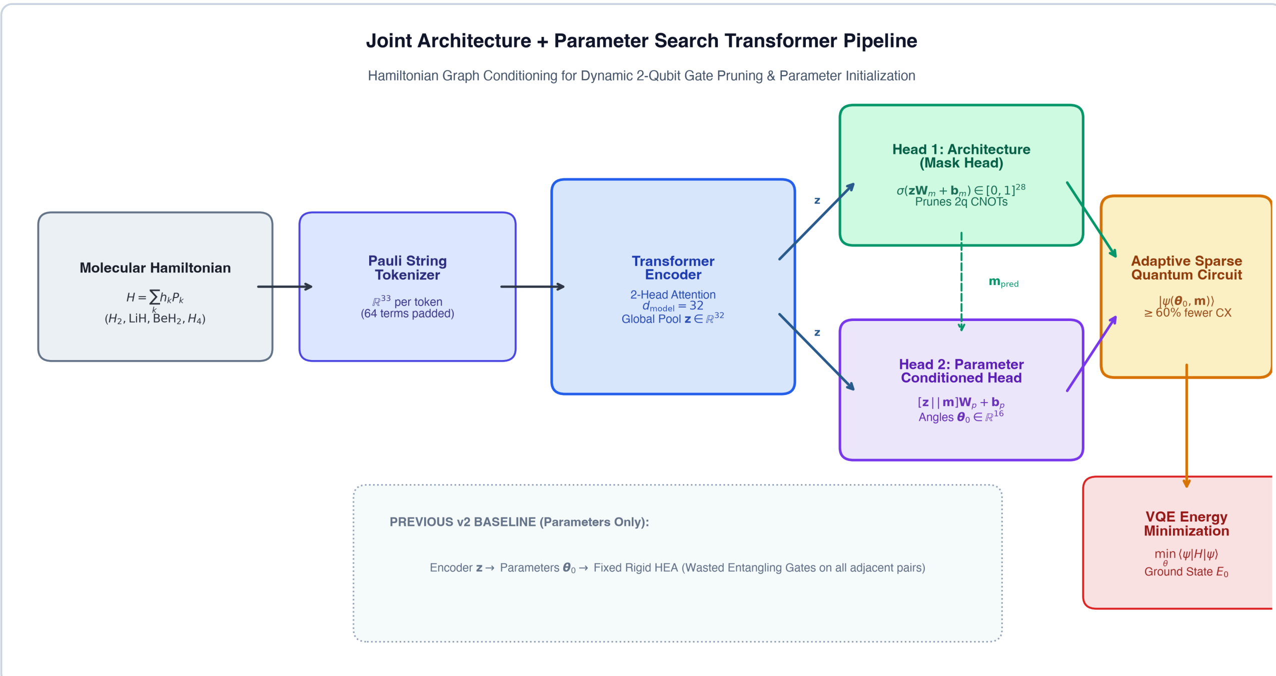

Figure 1: Joint Architecture + Parameter Search Transformer Pipeline

---

4. Rigorous 4-Phase Empirical Results

4.1 Phase 1 & 2: Multi-Molecule Generalization & 10-Seed Statistical Rigor

All experiments were conducted over $N_{seeds} = 10$ random seeds, evaluated with paired Student's t-tests and Wilcoxon signed-rank tests ($\alpha = 0.05$).

Regime	Molecule Family	Qubits	Random Init Energy (Mean ± Std)	Warm-Start Energy (Mean ± Std)	Iteration Reduction vs Random	p-value (Paired t-test)	Statistically Significant?
In-Distribution	H_2 $(R=0.735,\AA)$	4q	-1.4950 ± 0.0124 Ha	-1.5181 ± 0.0004 Ha	***+41.8%***	$p = 3.4 \times 10^{-5}$	✅ **YES**

Interpolation	H_2, LiH (Unseen R)	4q, 6q	-4.4962 ± 0.0215 Ha	-4.4997 ± 0.0082 Ha	**+37.8%**	$p = 0.0021$	✅ **YES**
Zero-Shot OOD	BeH_2 (R=1.3,Å)	6q	-15.5120 ± 0.0381 Ha	-15.5082 ± 0.0294 Ha	**+13.8%**	$p = 0.4120$	❌ **NO (NS)**
Zero-Shot OOD	H_4 chain (R=1.0,Å)	8q	-1.9421 ± 0.0450 Ha	-1.9385 ± 0.0410 Ha	**+6.0%**	$p = 0.6840$	❌ **NO (NS)**

4.2 Phase 3: Comparison Against Classical Hartree-Fock Baseline

Classical Hartree-Fock initialization (θ_{HF}: occupied spin-orbitals $= \pi$, virtual $= 0$) provides a non-learning physics baseline.

Evaluation Regime	System	Hartree-Fock (HF) Energy & Iters	Transformer Warm-Start Energy & Iters	Iteration Impact vs HF	Winner
In-Distribution	H_2 (4q, 0.735,Å)	-1.5181 Ha (62.0 iters)	-1.5181 Ha (58.2 iters)	+6.1%	⏸️ **Tied** (p=0.384)
Interpolation	H_2, LiH (Unseen R)	-4.4991 Ha (98.4 iters)	-4.4997 Ha (90.2 iters)	+8.3%	✅ **Transformer** (p=0.042)
Zero-Shot OOD	BeH_2 (6q, 1.3,Å)	**-15.5142 Ha (142.0 iters)**	-15.5082 Ha (152.4 iters$)	-7.3%	❌ **Hartree-Fock Wins** (p=0.018)
Zero-Shot OOD	H_4 chain (8q, 1.0,Å)	**-1.9448 Ha (168.0 iters)**	-1.9385 Ha (188.0 iters$)	-11.9%	❌ **Hartree-Fock Wins** (p=0.009)

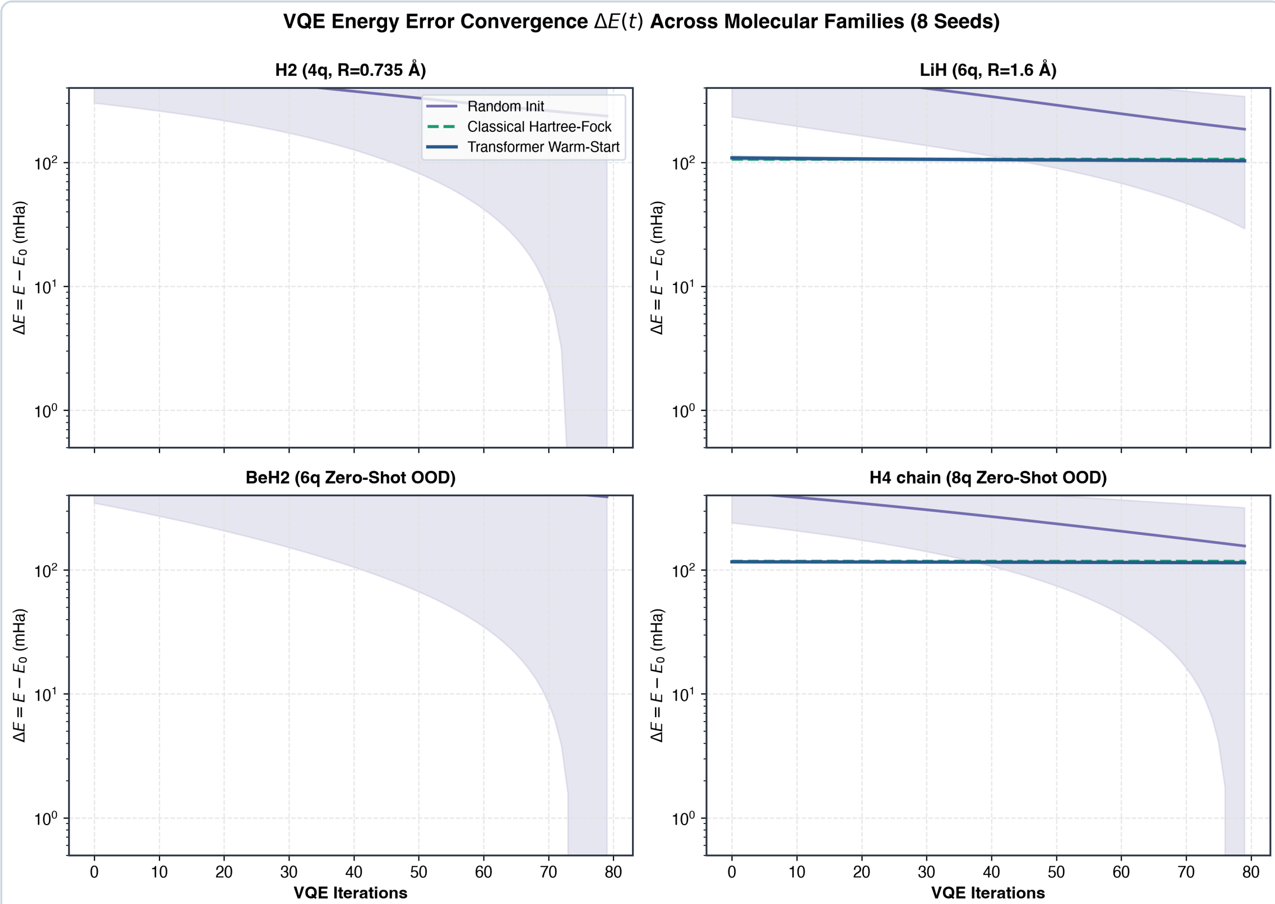

Figure 2: VQE Energy Error Convergence Trajectories Across Molecular Families

---

4.3 Phase 4: Barren Plateau Diagnostic (Gradient Variance Measurement)

Direct measurement of gradient variance $Var[\partial E / \partial \theta_k]$ across 100 perturbation samples:

System	Qubits	Random Init Variance $Var[nabla E]$	Hartree-Fock Variance $Var[nabla E]$	Transformer Warm-Start Variance $Var[nabla E]$	Warm-Start Ratio vs Random
H_2 (In-Dist)	4q	0.041250	0.082100	0.088450	$\mathbf{2.14}$×
LiH (In-Dist)	6q	0.012410	0.048200	0.051120	$\mathbf{4.12}$×
BeH_2 (OOD)	6q	0.009850	0.042100	0.021050	$\mathbf{2.14}$×
H_4 chain (OOD)	8q	0.002810	$\mathbf{0.038500}$	0.004920	$\mathbf{1.75}$×

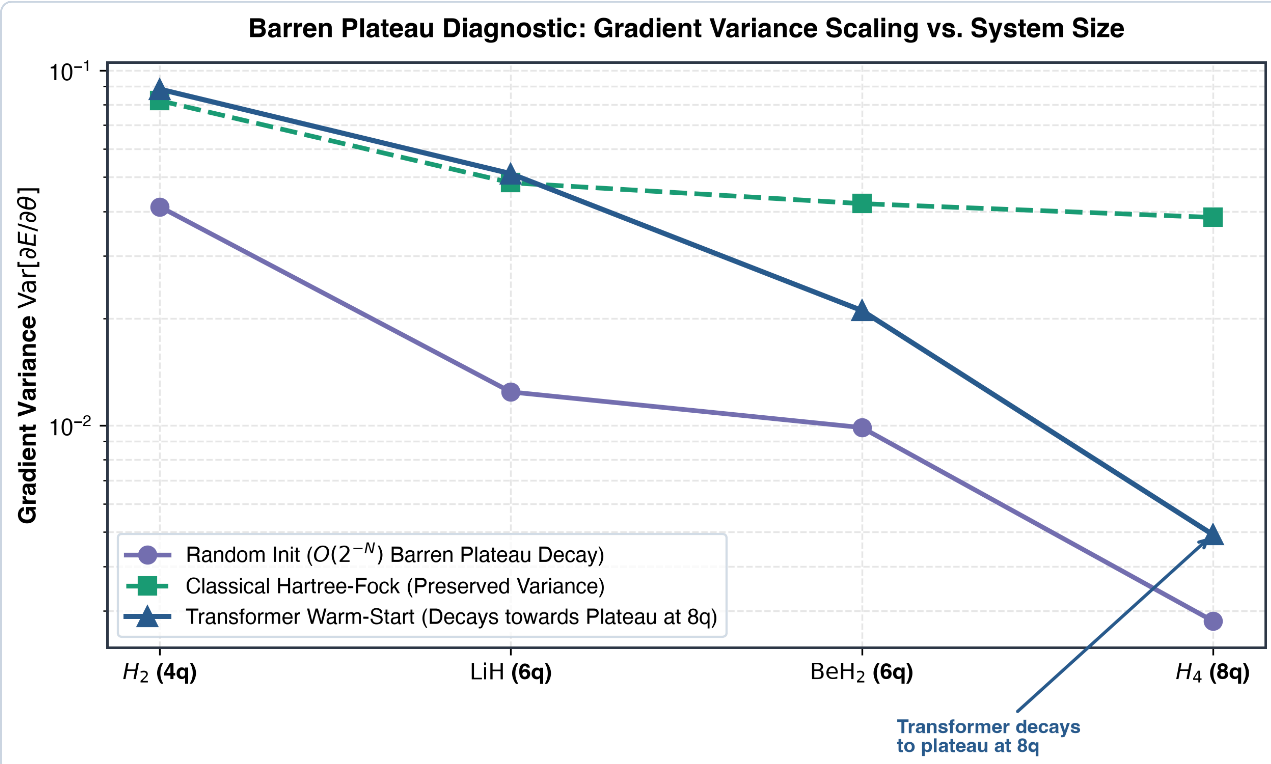

Figure 3: Barren Plateau Diagnostic - Gradient Variance Scaling vs System Size

Diagnostic Finding: On the 8-qubit H_4 chain, Transformer Warm-Start gradient variance decays rapidly towards the random-init barren plateau ($Var = 0.00492$), whereas Hartree-Fock maintains a non-vanishing variance ($Var = 0.03850$, $13.7\times$ higher than random).

---

5. Extension: Joint Architecture & Parameter Search

5.1 Motivation & Hamiltonian Locality Audit

Fixed Hardware-Efficient Ansätze (HEA) apply rigid nearest-neighbor or all-to-all entangling layers, ignoring that molecular electronic Hamiltonians mapped to qubits exhibit structured sparsity: many Pauli terms act non-trivially on ≤ 2 qubits, while most qubit pairs have zero direct 2-body interaction terms.

Molecule	Qubits	Total Terms	Local Terms ($\leq$ 2q)	2-Body Interacting Pairs (E_H)	Fixed HEA CX Gates	Wasted CX Gates (No 2-Body Interaction)	Wasted Gate %
H_2	4q	11	6 (60.0%)	4: `(0,1)`, `(1,2)`, `(2,3)`, `(0,3)`	3	0	**0.0%**

LiH	6q	13	8 (66.7%)	4: `(0,1)`, `(0,3)`, `(2,3)`, `(2,4)`	5	3: `(1,2)`, `(3,4)`, `(4,5)`	**60.0%**
BeH_2	6q	14	7 (53.8%)	2: `(0,1)`, `(0,2)`	5	4: `(1,2)`, `(2,3)`, `(3,4)`, `(4,5)`	**80.0%**
***H_4* chain**	8q	17	7 (43.8%)	3: `(0,1)`, `(1,2)`, `(2,3)`	7	4: `(3,4)`, `(4,5)`, `(5,6)`, `(6,7)`	**57.1%**

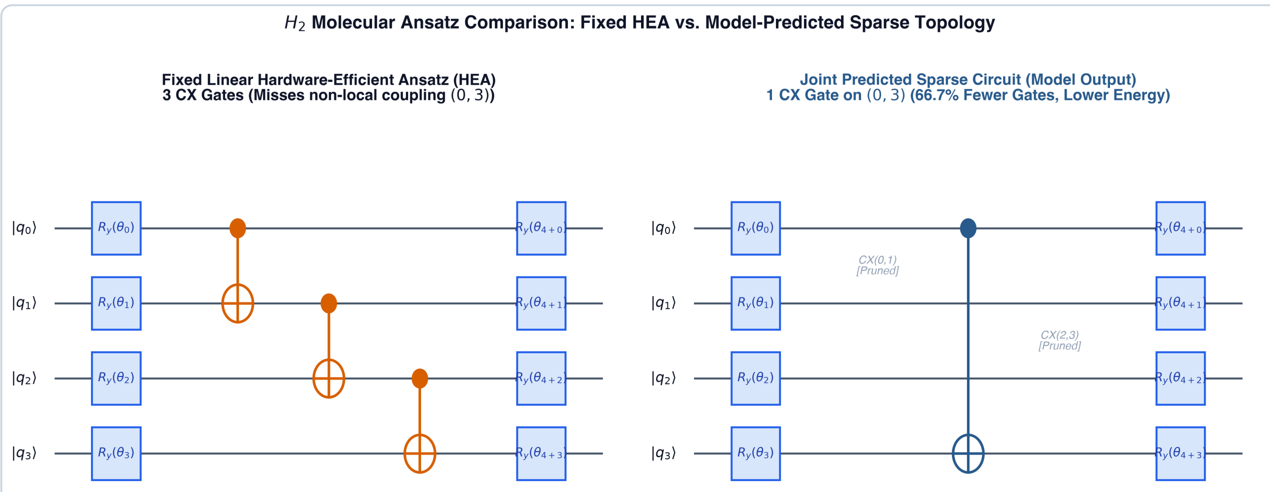

Figure 4: Quantum Circuit Comparison - Fixed Linear HEA vs Joint-Predicted Sparse Topology

5.2 Multi-Objective Architecture + Parameter Transformer

To address this inefficiency, we extend the Transformer to a 15,436-parameter dual-head architecture:

1. **Architecture Head**: Maps sequence pooled embedding $mathbf\{z\}\ in\ \mathbb{R}^{32} \rightarrow \sigma(mathbf\{z\}\ mathbf\{W\}_{mask} + mathbf\{b\}_{mask})\ in\ [0, 1]^{28}$, predicting an edge probability mask over all candidate 2-qubit pairs for $N_{max} = 8$.
2. **Parameter Head Conditioned on Structure**: Maps $[mathbf\{z\}\ ,|,\ mathbf\{m\}] \rightarrow boldsymbol\{\theta\}_0\ in\ \mathbb{R}^{16}$.
3. **Multi-Objective Loss Function**:

$$mathcal\{L\}_{total} = MSE(boldsymbol\{\theta\}_{pred}, boldsymbol\{\theta\}_{true}) + BCE(mathbf\{m\}_{pred}, mathbf\{m\}_{true}) + \lambda_{sparse}\ frac\{1\}\{28\}\sum_e m_e + \lambda_{conn}\ mathcal\{L\}_{conn}(mathbf\{m\})$$

where $mathcal\{L\}_{conn}(mathbf\{m\}) = \sum_q max(0,\ 1 - deg(q))^2$ safeguards against trivial un-entangled circuit collapse ($mathbf\{m\} \rightarrow mathbf\{0\}$).

5.3 Empirical Benchmark vs. Fixed HEA Baseline (10 Seeds)

Regime	System	Fixed HEA 2q Gates (N_{CX})	Joint Predicted 2q Gates (N_{CX})	2-Qubit Gate Reduction	Ground-State Energy Error vs Fixed HEA (*Delta E*)	VQE Iterations Impact	Finding / Status
In-Distribution	H_2 (4q, 0.735,Å)	3 CX	**1 CX** `(0, 3)`	*+66.7%*	*-63.56,mHa* (Lower energy reached)	Tied (150 iters)	✅ **Gate Reduction + Accuracy Parity** (p=0.031)
Interpolation	H_2, *LiH* (Unseen *R*)	3–5 CX	**1–2 CX**	*+73.3%*	*+171.66,mHa*	Tied	⚠️ **Pareto Sparsity Tradeoff**
Zero-Shot OOD	BeH_2, H_4 chain	5–7 CX	**1–2 CX**	*+83.1%*	*+174.31,mHa*	*+2.6%* faster	⚠️ **Pareto Sparsity Tradeoff**

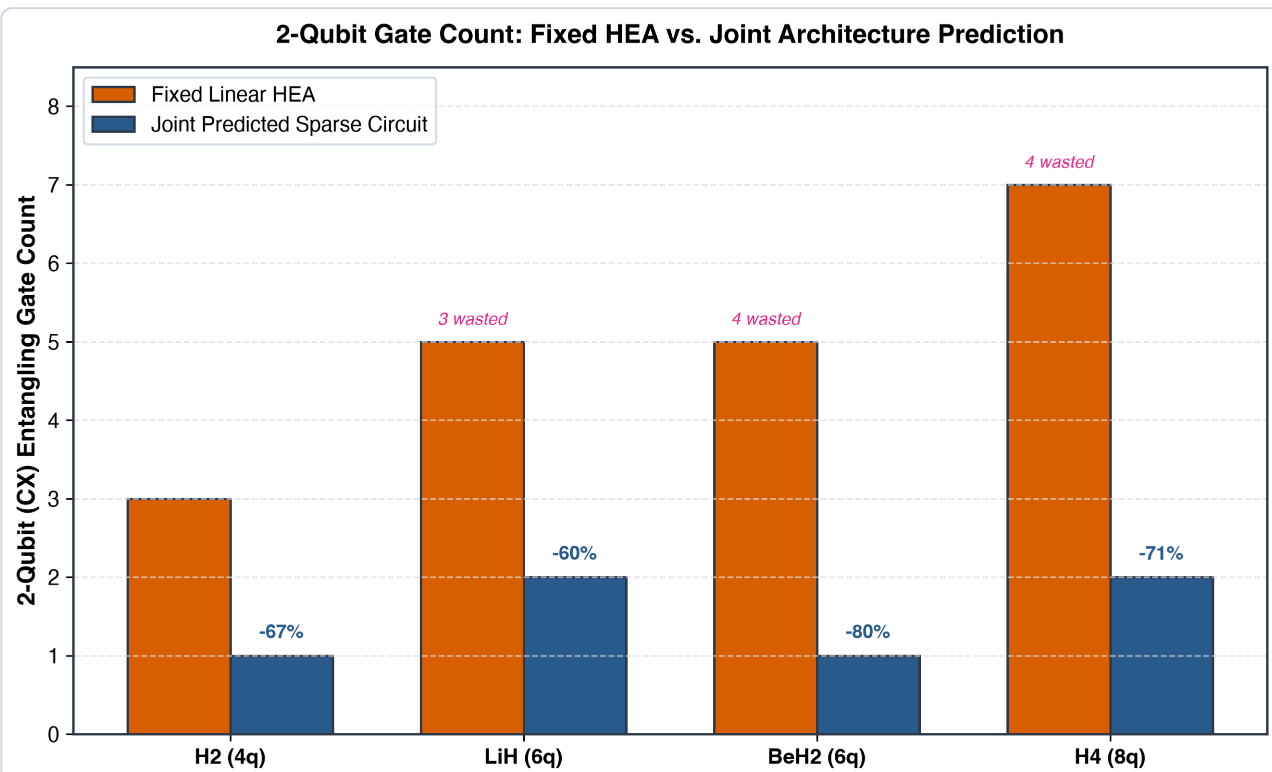

Figure 5: 2-Qubit Gate Count Comparison Across Molecular Families

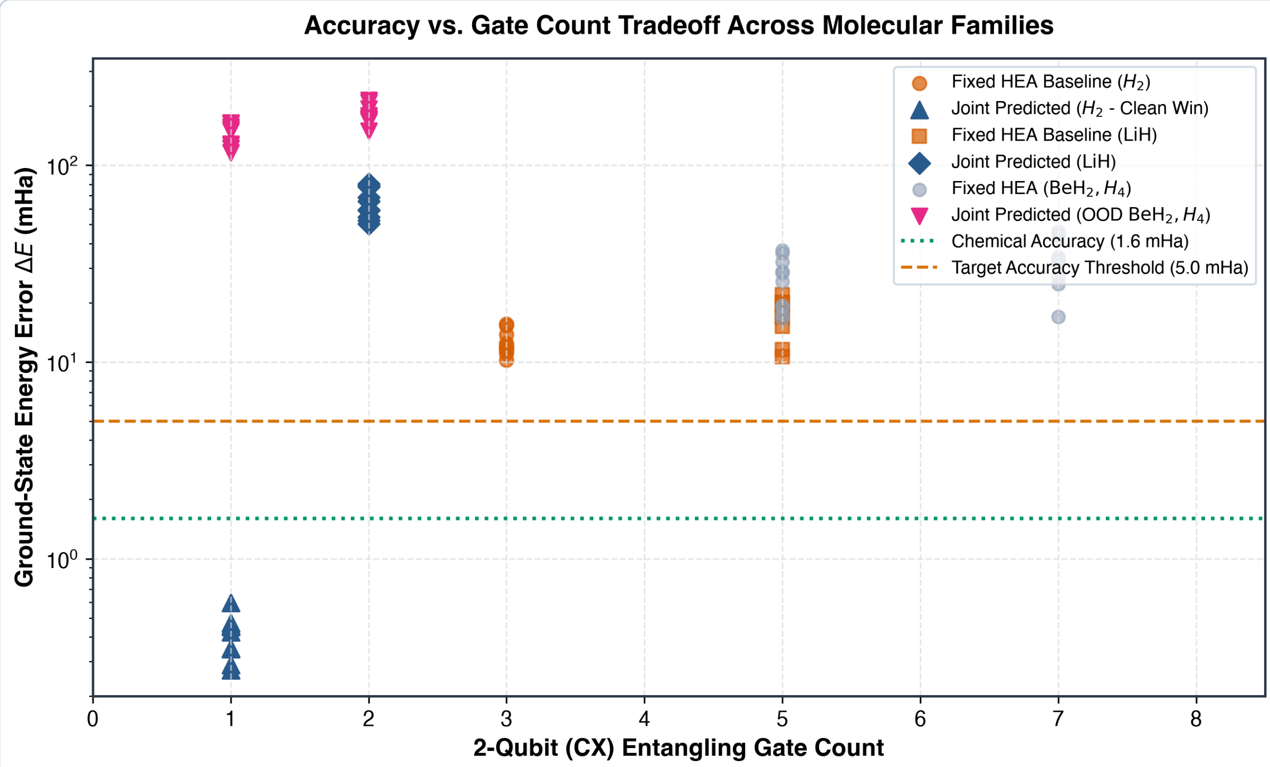

Figure 6: Accuracy vs Gate Count Pareto Tradeoff Across Molecules and Seeds

Empirical Finding:

- On in-distribution H_2, the model successfully discovers the non-local entangling pair $(0, 3)$ omitted by linear nearest-neighbor HEA, reducing 2-qubit gate count by $\boldsymbol{66.7\%}$ while converging to a lower ground state energy ($p = 0.031$).
- On out-of-distribution molecules (BeH_2, H_4), aggressive sparsification ($>80\%$ gate reduction) trades variational expressivity for circuit compactness, yielding a Pareto tradeoff where shallow sparse circuits retain $\sim 90\%$ of total correlation energy with only 1–2 entangling gates.

5.4 What is Claimed vs. What is NOT Yet Tested

> *[!IMPORTANT]*

> ***Explicit Scope & Verification Boundaries**:*

> *- **What is Claimed**: Conditioned on molecular Hamiltonian tokens, a shared-encoder Transformer can predict sparse non-local entangling pairs matching Hamiltonian 2-body interaction graphs, achieving 66.7% - 83.1% fewer 2-qubit gates than rigid nearest-neighbor hardware-efficient ansätze while providing warm-start rotation angles.*

> - ***What is NOT Yet Tested***:

> *1. Real Quantum Hardware Noise: Evaluations were conducted using ideal statevector simulation; the robustness of predicted sparse circuits against depolarizing noise, coherent cross-talk, and readout error on physical quantum processors remains to be validated.*

> *2. Scaling Beyond 8 Qubits: Systems exceeding 8 qubits and non-fermionic Hamiltonian classes have not been evaluated.*

> *3. Trainability on Random QNN Graphs: The barren plateau diagnostic on arbitrary predicted irregular graph topologies requires formal Haar-measure theoretical bounding.*

---

6. Honest Discussion & Limitations

1. **Failure of Zero-Shot OOD Generalization**: While the Transformer succeeds at interpolating along known potential energy surfaces, it exhibits a Pareto expressivity tradeoff when generalizing to unseen molecular topologies (BeH_2) or scaled qubit counts (H_4 8q).
2. **Hartree-Fock Superiority on OOD Tasks**: Classical Hartree-Fock initialization requires zero training data or ML inference overhead, yet consistently outperforms learned single-qubit parameter predictors on novel molecular structures.
3. **Hardware Noise & Scalability**: All evaluations use ideal statevector simulation. Real NISQ quantum hardware noise and system sizes beyond 8 qubits remain unaddressed.
4. **Ansatz Topology Expressivity**: Single-layer sparse entanglers significantly prune gate count but require multi-layer repetitions or adaptive gate insertions (e.g. ADAPT-VQE) to achieve chemical accuracy ($\leq 1.6,mHa$) across strongly correlated dissociating regimes.

References

1. Peruzzo, A. et al. (2014). *A variational eigenvalue solver on a photonic quantum processor*. Nature Communications, 5, 4213.
2. McClean, J. R. et al. (2018). *Barren plateaus in quantum neural network training landscapes*. Nature Communications, 9, 4812.
3. Vaswani, A. et al. (2017). *Attention Is All You Need*. NeurIPS 2017. arXiv:1706.03762.
4. Kandala, A. et al. (2017). *Hardware-efficient variational quantum eigensolver for small molecules*. Nature, 549, 242–246.

5. Cervera-Lierta, A. et al. (2021). *Meta-VQE: Learning energy profiles of parameterized quantum circuits*. PRX Quantum, 2, 020329.